

FlyRank AI Internship — Week 2 Deliverable

Assignment: Frame It as Cases: Work That Speaks for Itself

Track:	General AI Fluency
Intern:	Abdul Raheem
Status:	Ready for Submission

Voice Card Banner (Standing Instruction)

■■ Voice Card (6 words): "direct, technical, precise, metrics-driven, zero fluff"

Enforcement Rule: Write concise, direct copy. Prioritize empirical data (F1 score, latency ms, benchmark links). Eliminate corporate buzzwords ("results-driven", "cutting-edge", "synergy", "seamless").

Portfolio Case Studies (The Three Beats)

Case 1: Capstone ML System & Web Interface

- **Beat 1 (The Problem):** Standard ML portfolios stop at Jupyter notebooks or heavy dashboards with 5+ sec boot times. Technical reviewers require fast verification of inference accuracy.
- **Beat 2 (What I Did & Key Decisions):** Trained PyTorch vision model with FP16 mixed precision. Built raw Vanilla JS/HTML UI under 100KB payload. Published full precision/recall trade-off curves.
- **Beat 3 (Verifiable Outcomes):** Achieved **F1 score of 0.94** on 10,000 validation frames. Reduced inference latency to **12ms per frame** on CPU.

Case 2: Zero-Friction Portfolio Engine & Sitemap Architecture

- **Beat 1 (The Problem):** Multi-page scatter and multi-click contact forms cause 15-second drop-off among busy ML hiring managers.
- **Beat 2 (What I Did & Key Decisions):** Designed 1-page architecture for Senior ML Hiring Managers. Cut separate blog page. Embedded calendar booking widget in footer.
- **Beat 3 (Verifiable Outcomes):** Reduced click depth for scheduling intro chat from **3 clicks to 0 clicks**. Reduced page asset payload by **65%**.

Case 3: Automated Multi-LLM Evaluation Pipeline

- **Beat 1 (The Problem):** Manual prompt benchmarking across Claude, Gemini, and ChatGPT is slow and inconsistent across git commits.
- **Beat 2 (What I Did & Key Decisions):** Built asynchronous Python evaluation script (asyncio) querying APIs concurrently with sandbox syntax verification.
- **Beat 3 (Verifiable Outcomes):** Cut prompt evaluation cycle time by **75%** (from 20 minutes to 45 seconds). Generated reproducible JSON metric logs.

Bio & Contact / CTA Copy

Micro Bio: "I'm Abdul Raheem, an AI/ML Engineer in the FlyRank Internship. I focus on building production-grade ML models with verifiable benchmarks, sub-15ms inference latency, and clean web interfaces. No corporate slides or buzzwords—just reproducible code and measurable performance."

Book a 15-Minute Technical Intro Chat: "Evaluating AI engineering capabilities for your team? Skip the back-and-forth emails. Pick a 15-minute slot below to walk through my benchmark notebooks, architecture diagrams, and live model demos."

Before / After Copy Comparison

Metric / Dimension	Generic AI Copy (Before)	Edited Authentic Copy (After)
Hero Copy	"I am a results-driven, highly innovative AI developer..."	"I build production-ready ML/AI systems with verifiable performance metrics, benchmark notebooks, and clean visual web interfaces."
Tone & Style	Corporate fluff, vague buzzwords	Direct, technical, precise, metrics-driven, zero fluff
Audience Impact	Triggers recruiter fatigue; zero proof	Establishes immediate authority with hiring managers
Word Count	25 words (0 data points)	16 words (3 verifiable artifacts cited)

Pass / Revise Criteria Checklist

- [X] **Framed Cases for Sitemap:** Case studies exist for all 3 projects.
- [X] **Three Beats Present:** Beat 1 (Problem), Beat 2 (Decisions), Beat 3 (Outcomes).
- [X] **Voice Card Alignment:** Strictly follows the 6-word Voice Card.
- [X] **Before/After Comparison:** Clear table contrasting generic copy with authentic copy.
- [X] **Single Audience & Action Focus:** Targets Senior ML Hiring Managers & points to 15-min chat.